Machine Learning

Defect Prediction on Production Line

Team:
AIT BACHIR Romuald
ALLEMAND Fabien
FLORET Arthur
JARDOT Charles

Abstract

The primary goal in detecting anomalies using machine learning in the context of a industry is to develop a model able to detect whether an item is defective or not using data related to said item. This model would be a major improvement for the industry since it takes time to manually test every single part.

As expected from production lines' data, defective and non-defective pieces are not equally represented. Since the imbalance between classes harm the model's learning process, balancing them is one way to deal with this issue.

The Naïve Bayes classifier was chosen to answer classify items on the production line because it was evaluated as the most performant model at this task among various other machine learning models. The results demonstrate that such a simple algorithm can already predict defects, which is encouraging for the future when more data is collected and heavier models are implemented.

Contents

List of Figures

1 Introduction

The goal of this project is to develop an AI based solution for a real world industry problem faced by Valeo [1].

Valeo [2] is an industry leading French automotive supplier. In order to stay competitive, the company wants to develop a system that is able to identify defects on products before testing.

Four files containing relevant data are available. The data are mainly values measured during production on different mounting stations as well as additional measures performed on test benches.

The target is to find the best prediction: Output = f(inputs)

2 Data Analysis

The goal of this section is to take a first look a the data in order to understand what it represents, then proceed to an in-depth analysis of said data.

2.1 Data Description

Data is contained in four different comma-separated values (csv) files: inputs and output for both training and testing. The very first line of each file correspond to the headers. The following lines contain informations about a given product where each product is identified by a unique code:

ID = PROC_TRACEINFO = It's a unique code given to the product.
Example: I-B-XA1207672-190701-00494.

- XA1207672 is the reference.

- 190701 is the date: here 01st of July of year 2019.

- 00494 is the unique code given to the product, whatever it happens, the product will have this id number frozen forever.

This number is increased by 1 each time we process a new product, every 12 seconds. So for example: I-B-XA1207672-190701-00495 is the next product.

Input features are measures (numerical values) collected on different assembly stations with the sensors or devices connected to Programmable Logic Controllers which are storing all of them to keep the full quality traceability. Examples : OP070_V_1_angle_value, OP120_Rodage_I_value...

This is the result value of OP130 (test bench). Value 0 is assigned to OK samples (passed) and value 1 is assigned to KO samples (failed). This is the combined result of multiple electrical, acoustic and vibro-acoustic tests.

2.2 In-Depth Analysis

Let us put the testing data aside for now and take a closer look a the training data.

The training data consists of 13 measures (inputs features) made during production line on just shy of 35000 products. This measures correpond to various forces, angles, screwing torques... By taking a look at the distribution for each feature (Figure 1) we can clearly see that every one of them has a different shape.

A quick glance at the pairwise correlation matrix reveals that there are no obvious correlation between any features (Figure 2).

Looking at the results of a PCA (Figure 3 4) or at multiple histograms (with items grouped by class for each feature, as displayed in Figure 5), there is no feature that clearly helps separating the two classes.

An in-depth look at the box-plots (Figure 6) reveals that there might be some outliers in the dataset.

École d'ingénieurs
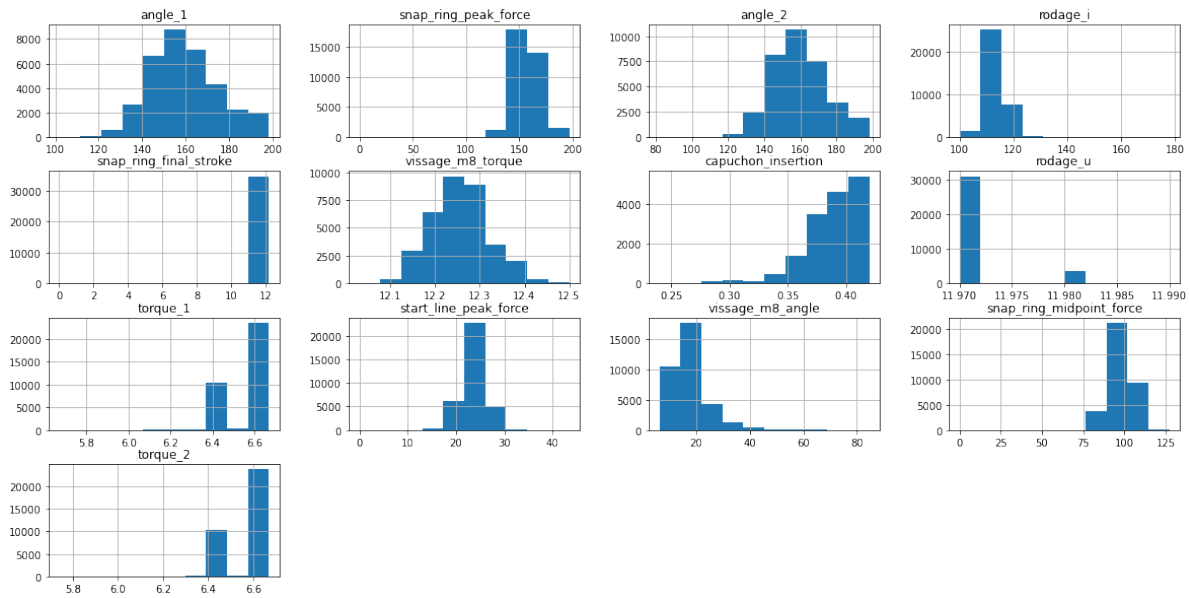Télécom Physique Strasbourg

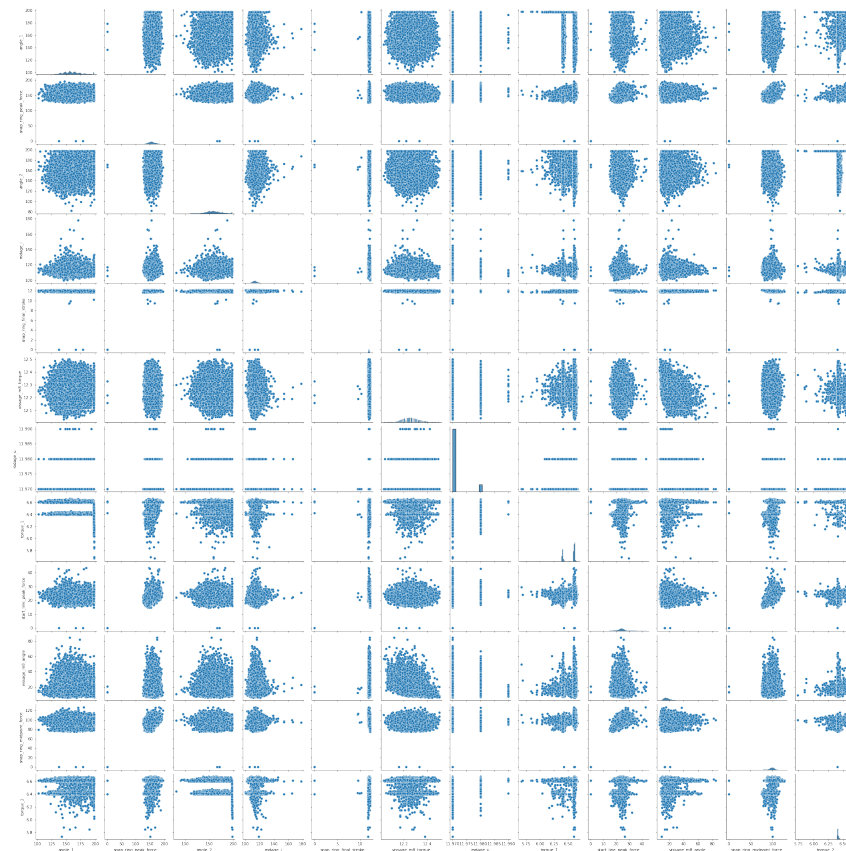Figure 1: Distributions of the input features

Figure 2: Features pairwise correlation matrix

Figure 3: PCA results (projection of the training population on the three most relevant dimensions)

Figure 4: PCA results (projection of the training population on the two most relevant dimensions)

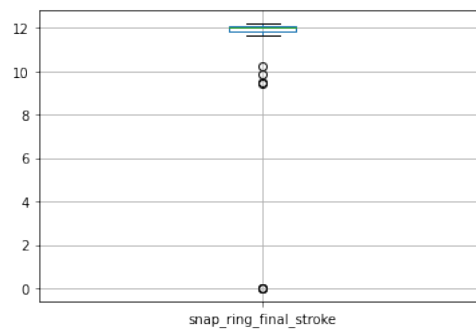Figure 5: Distribution of the *angle_2* feature with items grouped by class

Figure 6: Box-plot of the *snap_ring_final_stroke* feature

One feature requires more attention. *capuchon_insertion* is the only feature containing missing values. In fact, more than 50% of the sample population are missing this value.

The training data also contains the result value of the test bench for each product tested on the production line. As expected, there are only a few of defective items. Although this is good for the company, it means that **less than 1%** of the training population are defective items. The effectives of the two classes are very unbalanced which could really impact our models.

3   Data Preparation

The data analysis performed in the previous section revealed some flaws in the dataset. The two most noticeable are the presence of outliers and the population imbalance between the two classes. These two flaws are known for harming the learning of most ML models. Different attemps at solving these issues are presented in this section.

3.1   Outliers

3.2   Balance Classes

The following paragraphs deal with class balancing, explaining the importance of balanced classes and three different approaches two rectify the imbalance.

As said in the previous section, a mere 0.008% of the population is defective individuals. In other words, the class of individuals which needs to be detected is truly under-represented in the dataset. Without any pre-processing such a dataset is very unlikely to give good results with most of the ML models. As

Figure 7: Confusion matrix of a k-Nearest Neighbors classifier with the base dataset


Figure 8: Confusion matrix of a k-Nearest Neighbors classifier with the dataset balanced by removing individuals

shown in Figure 7, there are not enough defective individuals for the classifier to properly learn how to distinguish a defective individual from a valid individual. As a result, the classifier simply over-fit the valid class.

Remarque | Over-fitting the over-represented class is very hazardous as the accuracy of the classifier will remain really high eventhough the classifier is not classifying anything. Hopefully the confusion matrix shows the classification details.

The easiest method to deal with unbalanced classes is to remove valid individuals from the training dataset. This process is fairly straight forward (valid individuals can be randomly selected in the dataset) and does not require a lot of processing (selected individuals are simply deleted).
Figure 8 clearly shows that the classifier is not over-fitting the valid class anymore. In that sense, the results are better but they are still far from what is expected from a performant classifier.

This method is very arguable for two reasons. First, deleting individuals obviously implies a loss of data. Eventhough the individuals are randomly selected, some important features can disappear. In other words, the remaining data is not necessarily representative of the global population. On top of that, the resulting dataset is smaller meaning the model will have less instances to learn on. In this problem, this will dramatically impact the results because there is already few data.

Of course it is also possible to work the other way around and artificially create individuals instead of deleting others. The easiest method is actually to duplicate defective items until there are as many as the valid ones. This compensate the imbalance between the two classes without knowing anything about the features values. Figure 9 shows the effect on the classifier is significant.
Another way that has been implemented during this project is to slightly modify individuals during the duplication. This way new individuals are actually new individuals (not only duplicates) adding information to the overall dataset. This method assumes individuals from the same class are close to one another on

Figure 9: Confusion matrix of a k-Nearest Neighbors classifier with the dataset balanced by duplicating individuals

every dimension (i.e. regarding every feature) and changing a feature's value by a small amount will not induce a change of class. The way it has been implemented is surely not ideal but rather intuitive. A random amount of features are selected during each "duplication" and the value is modified by a very small percentage (typically 0.005).

The final method in order to balance classes consists in leaving the classes inbalanced in the dataset but during computation attributing a more significant weight to the individuals from the under-represented class. Technically speaking this is equivalent to duplicating individuals but it requires less memory, it is easier for computation and is already implemented for some algorithms from well known ML libraries.

4   Model Selection

Choosing the ML model that works the best for a given problem is a key factor for good results but it can be very time consuming. The following pages present the results obtained with several ML algorithms. Each method will be briefly introduced and tested using the training dataset. The resulting classifier are evaluated using three different tools:

- Confusion matrix

- Precision, recall and F1 score

- ROC AUC Score

These three tools are presented and explained in Section 4.1.
Knowing the data is labeled all the methods but the last one are based on supervised learning.
Moreover it has already been demonstrated in Section 3.2 that balancing classes dramatically improves results so the ML models are tested with an augmented training dataset containing duplicated individuals.

4.1   Evaluation Method

There are many tools to analyse the performances of ML models. Three tools will be thoroughly used later. The goal of the following paragraphs is to give a quick understanding at what they represent and how to use them properly.

The confusion matrix is a tool specifically used for classification problems. This matrix simply represents the results of the classification made by a given classifier. For instance in Figure 9, the classifier classified 237 valid items as defective (top right).
By looking at a confusion matrix, it is fairly easy to see what the classifier did wrong by comparing it to the ideal matrix which would contains zeros for the false positive (top right) and false negative (bottom left)

emplacements.

The values contained in the confusion matrix also allow to compute the precision, recall and f1 score of a classifier:

- Precision: how accurate is the classifier when classifying a defective item (closer to one is better)

- Recall: how good is the classifier at finding defective items in the population (closer to one is better)

- f1 score: the harmonic mean of precision and recall (closer to one is better)

Precison and recall are linked values meaning there is always a tradeoff: if precison rises, recall tends to decrease. So it is important to adjust precision and recall according to the application. In this problem, a company would most likely prefer to double check an item rather than selling a defective one so a strong recall would be ideal, at the expence of precision.

Finaly, Receiver Operating Characteristic (ROC) curve plots the true positive rate against the false positve rate (i.e. sensitivity against specificity, Figure 10). It allows to distinguish the classifier from a random classifier. It is generally associated to the ROC Area Under the Curve (AUC) score. A ROC AUC score over 0.5 means the classifier is better than a random classifier. An ideal classifier would show a ROC curve following the top left edge of the graph and a ROC AUC score of one.

4.2   Naive Bayes Classifier

Naive Bayes classifiers use Bayes'theorem making the "naive" assumption that every pair of features are independant given the value of the class variable [4] (i.e. $P(x_i|y, x_1, ..., x_{i-1}, x_{i+1}, ..., x_n) = P(x_i|y)$) leading to the following equality:

$$P(y|x_1, ..., x_n) = \frac{P(y)P(x_1, ..., x_n|y)}{P(x_1, ..., x_n)} = \frac{P(y)\prod_{i=1}^{n} P(x_i|y)}{P(x_1, ..., x_n)}$$

The following results correspond to a Gaussian Naive Bayes [7] classifier that is very common when dealing with numerical continuous values. However it also assumes that every features follows a Gaussian distribution.

The confusion matrix in Figure 10 shows poor results. This is expected because Figure **??** demonstrates some features do not follow a Gaussian distribution. In addition, the assumption that features are independant from one another might not be relevant with this dataset. For instance, the *snap_ring_peak_force* might increase with the value of *anlge_1*.

4.3   k-Nearest Neighbors

The k-nearest neighbors algorithm works by classifiying an individual by looking at its neighbors [0]. The classifier simply look at the k nearest neighbors of the individual in the training population and attributes the class that correspond the best to the neighbors (the neighbors vote for their class, each vote is weighted either evenly accros the neighbors or according to the distance between the individual and the voting neighbor [8]).

According to Figure 11, this algorithm creates a model that is close to a random classifier. Using cross-validation, the average precision is lower than 1% and the average recall is lower than 2% (with very little deviation in each case). This means that, no matter the training set, the classifier is no able to distinguish a defective item from a correct one.

4.4   Random Forest Classifier

Random Forest Classifier works by creating multiple Decision Trees on various subsets of the data, the decision of the forest for classification is simply the class that has been predicted by most of the trees [0]. A Decision
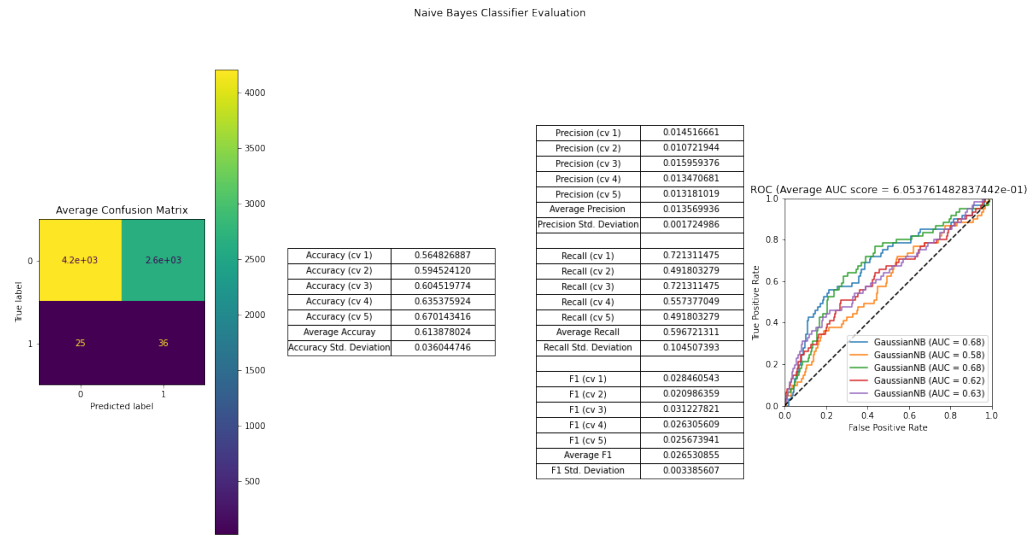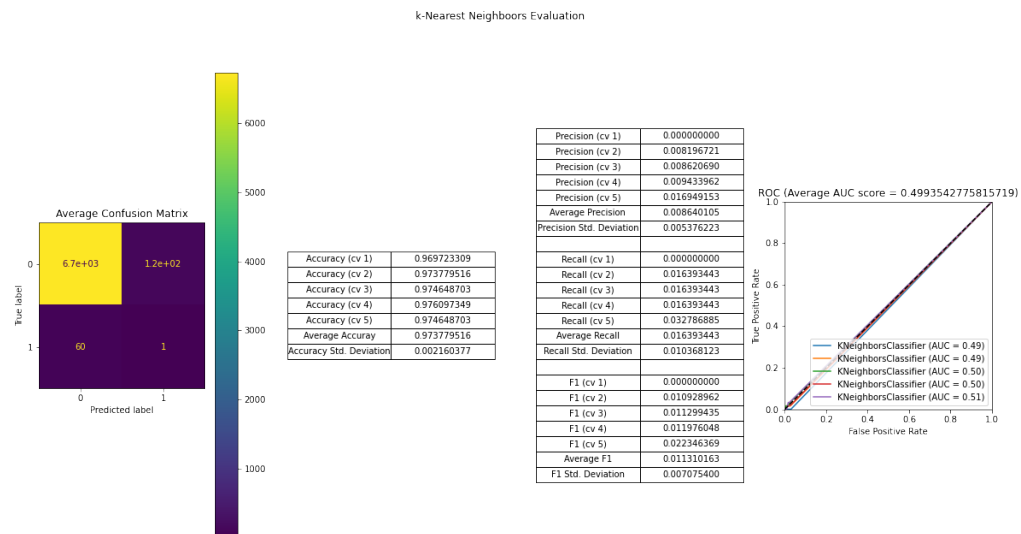
Figure 10: Naive Bayes Classifier Evaluation

Figure 11: k-Nearest Neighbors Evaluation

Random Forest Classifier Evaluation

Figure 12: Random Forest Classifier Evaluation (with *class_weight* set to "balanced")

Tree is built by finding the rule which devides the best the population of each node, this can be the rule that isolate most individuals from the same class (this method is called information gain and is based on the concept of entropy)[0]. A node cannot be devided when there are not enough individuals (in the node or in the to-be-created child-node), then the leaf is labeled with the class with the most representants on the node.

Random Forest Classifiers have a *class_weight* attribute in Scikit-Learn in order to address imbalance between classes[6]. However it seemed to perform worse than training the forest with the dataset containing duplicates (Figure 12). The ROC AUC score on Figure 13 shows just-above-random results. But looking at the precision and recall shed the light on a useless model. Not being able to detect any defective item makes it a very bad candidate.

4.5 Multilayer Perceptron

Multilayer Perceptron is part of the non-linear supervised learning branch. MLPs are characterised by iterative learning on the perceptrons present in the network layers [9]. Several arguments can be adjusted in order to modify the shape and parameters of the MLP:

- hidden_layer_sizes : number of neurons in each hidden layer

- activation : activation function for the hidden layer

- solver : solver for weight optimization

- alpha : strength of the L2 regularization term

- batch_size : size of minibatches for stochastic optimizers

- learning_rate : learning rate schedule for weight updates

- max_iter : maximum number of iterations

- random_state : random number generation for weights and bias initialization

Random Forest Classifier Evaluation

Figure 13: Random Forest Classifier Evaluation

- early_stopping : early stopping indicator (terminate training when validation score is not improving)

- n_iter_no_change : maximum number of epochs to not meet tolerance for the optimization improvement

The MLP that gave the best results is composed of an input layer with 12 perceptrons and an output layer with one perceptron.

It is important to note that in the sklearn module, it is not possible to modify the cost function, which by default is the LogLoss function. The Logloss function minimizes the inverse of the log likelihood:

$$LL = -\frac{1}{n}\sum_{i=0}^{n}\left(y_i \log(p_i) + (1 - y_i)\log(1 - p_i)\right)$$

The loss is calculated for each batch, n is the value of the batch size, y_i is the output of the neural network and p_i is the associated probability.

There is no method to anticipate the most relevant hyperparameters. The only way is to test several combinations by changing only one hyperparameter at a time and to observe the performance of the resulting network on our data. The best parameters obtained during testing are presented in Figure 14.

Figure 15 shows decieving results. Despite the time invested in training the network, it is not able to perform better than a random classifier on average (ROC AUC score below 0.5 on Figure 15). This is another bad candidate for the problem.

4.6 Novelty Detection

Anomaly or Novelty detection is a subset of machine learning particularly fitting to unbalanced class problems [5]. Knowing only the data and the estimated proportion of defects, the model will train itself and learn what a valid individual is supposed to look like [3]. It will then be able to classify new individuals: everything that is different than a valid item will be classified defective. Consequently, this means that the model is able to detect defects that are not represented in the training datset. This ability might be very be a advantge for this application as the model will stay relevant even when new defects occur. For instance, the tools on the assembly line will probably wear down leading to the apparition of new defects.

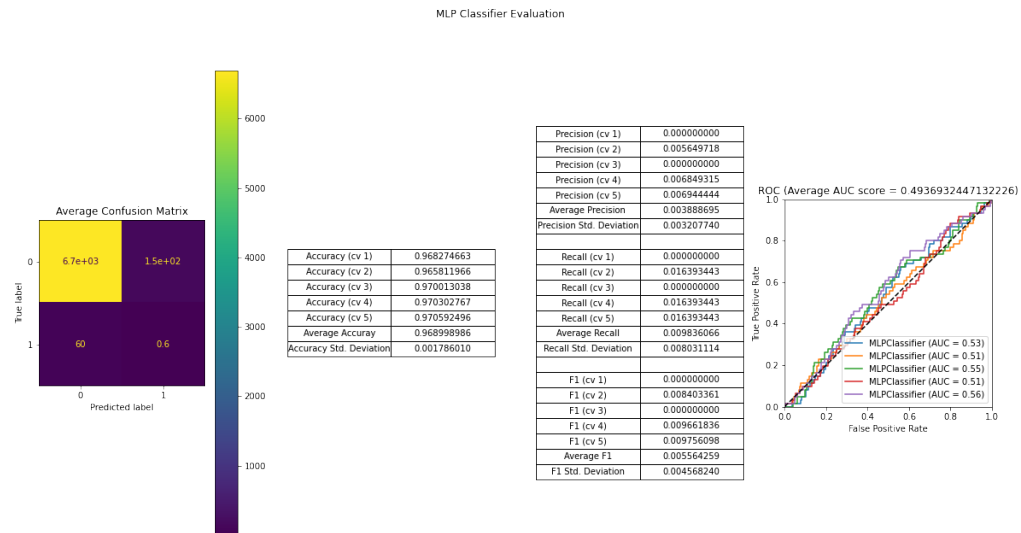| Hyperparameters	Selected value
hidden_layer_sizes	(200,150,100,50)
activation	hyperbolic tangent
solver	stochastic gradient descent
alpha	10^{-3}
batch_size	128
learning_rate	adaptive
max_iter	2000
random_state	1
early_stopping	True
n_iter_no_change	6

Figure 14: Best parameters for the MLP

Figure 15: MLP Classifier Evaluation

5   Model Fine Tuning

6   Conclusion

References

[1] Defect prediction on production line by valeo. `https://challengedata.ens.fr/challenges/36`. Accessed: 2023.

[2] Valeo. `https://www.valeo.com/fr/`. Accessed: 2023.

[3] Grab N Go Info. One class svm for anomaly detection — unsupervised machine learning. `https://www.youtube.com/watch?v=OIkFnHpUUjE&ab_channel=GrabNGoInfo`. Accessed: 2023.

[4] Scikit-Learn. Naive bayes. `https://scikit-learn.org/stable/modules/naive_bayes.html`. Accessed: 2023.

[5] Scikit-Learn. Novelty and outlier detection. `https://scikit-learn.org/stable/modules/outlier_detection.html`. Accessed: 2023.

[6] Scikit-Learn. sklearn.ensemble.randomforestclassifier. `https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html`. Accessed: 2023.

[7] Scikit-Learn. sklearn.naive_bayes.gaussiannb. `https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.GaussianNB.html`. Accessed: 2023.

[8] Scikit-Learn. sklearn.neighbors.kneighborsclassifier. `https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html`. Accessed: 2023.

[9] Scikit-Learn. $sklearn.neural_network.mlpclassifier.$. Accessed: 2023.

Wikipedia. Decision tree learning. `https://en.wikipedia.org/wiki/Decision_tree_learning`. Accessed: 2023.

Wikipedia. k-nearest neighbors algorithm. `https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm`. Accessed: 2023.

Wikipedia. Random forest. `https://en.wikipedia.org/wiki/Random_forest`. Accessed: 2023.